

9: Further Methods

John H Maindonald

March 16, 2018

```
doFigs <- FALSE
```

```
fig9.1 <- function(){  
  ## ---- smooth-ohms ----  
  ## Plot points  
  plot(ohms ~ juice, data=fruitohms)  
  ## Add smooth curve, using default  
  ## smoothing window  
  with(fruitohms,  
        lines(lowess(ohms ~ juice), col="gray", lwd=2))  
}
```

```
fig9.3 <- function(){  
  ## ---- ant111b ----  
  # ant111b is in DAAG  
  Site <- with(ant111b, reorder(site, harvwt,  
                                FUN=mean))  
  stripplot(Site ~ harvwt, data=ant111b,  
            scales=list(tck=0.5),  
            xlab="Harvest weight of corn")  
}
```

```
fig9.4 <- function(){  
  ## ---- Erie1 ----  
  plot(Erie, xlab="",  
        ylab="Level (m)")  
}
```

```

fig9.2A <- function(){
  if(!exists("eyeAmpM.gam"))
  eyeAmpM.gam <- mgcv::gam(amp ~ s(x,y), data=subset(eyeAmp, Sex=="m"))
  if(!exists("eyeAmpF.gam"))
  eyeAmpF.gam <- mgcv::gam(amp ~ s(x,y), data=subset(eyeAmp, Sex=="f"))
  lims <- range(c(predict(eyeAmpF.gam), predict(eyeAmpM.gam)))
  mgcv::vis.gam(eyeAmpM.gam, plot.type='contour', color="cm", zlim=lims, main="")
}
fig9.2B <- function(){
  if(!exists("eyeAmpM.gam"))
  eyeAmpM.gam <- mgcv::gam(amp ~ s(x,y), data=subset(eyeAmp, Sex=="m"))
  if(!exists("eyeAmpF.gam"))
  eyeAmpF.gam <- mgcv::gam(amp ~ s(x,y), data=subset(eyeAmp, Sex=="f"))
  lims <- range(c(predict(eyeAmpF.gam), predict(eyeAmpM.gam)))
  mgcv::vis.gam(eyeAmpF.gam, plot.type='contour', color="cm", zlim=lims, main="")
}
fig9.2 <- function(){
  print("Run the separate figures fig9.2A() and fig9.2B()")
}

```

```

fig9.5A <- function(){
  ## ---- lagErie ----
  ## Panel A
  lag.plot(Erie, lags=3,
           do.lines=FALSE,
           layout=c(1,3),
           main=" ")
  mtext(side=3, line=2, adj=0,
        "A: Lag plots, at lags of 1, 2 and 3")
}
fig9.5B <- function(){
  ## ---- acfErie ----
  ## Panel B
  acf(Erie, main="")
  mtext(side=3, line=1, adj=0, cex=1.25,
        "B: Autocorrelation function")
}
fig9.5 <- function(){
  "Run the separate functions fig9.5A() and fig9.5B()."
}

```

```

fig9.6 <- function(){
  ## ---- gamErie ----
  df <- data.frame(

```

```

height=as.vector(Erie),
year=time(Erie))
obj <- mgcv::gam(height ~ s(year), data=df)
plot(obj,
      shift=mean(df$height),
      residuals=TRUE, pch=1,
      xlab="",
      ylab="Height of lake")
}

```

```

fig9.7 <- function(){
## ---- arima-sim ----
for (i in 1:6){
ysim <- arima.sim(list(ar=0.85),
                  n=200)
df <- data.frame(x=1:200,
                 y=ysim)
df.gam <- mgcv::gam(y ~ s(x),
                  data=df)
plot(df.gam,
      ylab=paste("Sim", i),
      residuals=TRUE)
}
}

```

```

fig9.8 <- function(){
## ---- Erie-fcast ----
erie.ar <- ar(Erie, order.max=1)
fc <- forecast::forecast(erie.ar, h=15)
plot(fc, main="",
      ylab="Lake level (m)")
# 15 time points ahead
}

```

```

fig9.9 <- function(){
## ---- mdb-gam ----
mdbRain.gam <- mgcv::gam(mdbRain ~ s(Year) + s(SOI),
                        data=bomregions)
plot(mdbRain.gam, residuals=TRUE, se=2, pch=1,
      cex=0.5)
}

```

```

fig9.10 <- function(){
  ## ---- ar1sims ----
  opar <- par(mar=c(2.1, 3.1, 0.25, 1.1), mgp=c(2.25,0.5,0))
  mdbRain.gam <- mgcv::gam(mdbRain ~ s(Year) + s(SOI),
                          data=bomregions)
  n <- dim(bomregions)[1]
  acf(resid(mdbRain.gam), ylab="MDB series")
  for(i in 1:5)acf(rnorm(n), ylab=paste("Sim",i),
                  col="gray40")
  par(opar)
}

```

```

library(DAAGviz, quietly = TRUE)

```

```

if(!requireNamespace('gamclass', quietly=TRUE))print('gamclass must be installed')
fromDate <- as.Date("2006-01-01")
dfDay06 <- gamclass::eventCounts(data=gamclass::airAccs, dateCol="Date",
                                from= fromDate, by='1 day', prefix="num")
dfDay06$day <- julian(dfDay06$Date, origin=fromDate)
dfWeek06 <- gamclass::eventCounts(data=gamclass::airAccs, dateCol="Date",
                                from=fromDate,
                                by='1 week', prefix="num")
dfWeek06$day <- julian(dfWeek06$Date, origin=fromDate)
year <- seq(from=fromDate, to=max(dfDay06$Date), by="1 year")
atyear=julian(year, origin=fromDate)

```

```

fig9.11A <- function(){
  if(!exists('dfday06')){
    print(paste("Unable to plot graph.",
                "Ensure package 'gamclass' is installed.))
    return()
  }
  dfDay06.gam <-
    mgcv::gam(formula = num ~ s(day, k=200), family = quasipoisson,
              data = dfDay06)
  av <- mean(predict(dfDay06.gam))
  plot(dfDay06.gam, xaxt="n", shift=av, trans=exp, rug=FALSE, xlab="",
        ylab="Estimated rate per day")
  axis(1, at=atyear, labels=format(year, "%Y"))
  mtext(side=3, line=0.75, "A: Events per day, vs date", adj=0)
}
fig9.11B <- function(){
  dfWeek06.gam <- mgcv::gam(num~s(day, k=200), data=dfWeek06, family=quasipoisson)

```

```

av <- mean(predict(dfWeek06.gam))
plot(dfWeek06.gam, xaxt="n", shift=av, trans=exp, rug=FALSE, xlab="",
      ylab="Estimated rate per week")
axis(1, at=atyear, labels=format(year, "%Y"))
mtext(side=3, line=0.75, "B: Events per week, vs date", adj=0)
}
fig9.11 <- function(){
  "Run the separate functions fig9.10A() & fig9.10B()"
}

```

```

fig9.12 <- function(){
  ## ---- fgl-scores2D ----
  scores <- predict(fgl.llda)$x
  xyplot(scores[,2] ~ scores[,1], groups=fgl$type,
          xlab="Discriminant 1",
          ylab="Discriminant 2",
          aspect=1, scales=list(tck=0.4),
          auto.key=list(space="right"),
          par.settings=simpleTheme(alpha=0.6, pch=1:6))
}

```

```

## ---- bronchit-first3 ----
## ---- bronchit-rp ----
set.seed(47) # Reproduce tree shown
fig9.13 <- function(){
  ## ---- treefig ----
  opar <- par(mar=rep(1.1,4))
  if(!exists('b.rpart'))b.rpart <- rpart(rfac ~ cig+poll, data=bronchit)
  plot(b.rpart)
  text(b.rpart, xpd=TRUE)
  par(opar)
}

```

```

fig9.14 <- function(){
  ## ---- rf-x-bronchit ----
  opar <- par(mar=rep(2.1,4))
  set.seed(31) # Reproduce the trees shown
  num <- 1:nrow(bronchit)
  for(i in 1:12){
    useobs <- sample(num, replace=TRUE)
    dset <- bronchit[useobs, ]
    b.rpart <- rpart(rfac ~ cig+poll, data=dset,
                     control=rpart.control(maxdepth=2))
  }
}

```

```

plot(b.rpart, uniform=TRUE)
text(b.rpart, xpd=TRUE, cex=1.2)
}
par(opar)
}

```

```

fig9.15 <- function(){
  ## ---- proximity-plot ----
  parset <- simpleTheme(pch=1:2)
  if(!requireNamespace('randomForest'))
    return('Fig 9.15 requires randomForest package')
  bronchit.rf <- randomForest::randomForest(rfac ~ cig+poll,
                                           proximity=TRUE,
                                           data=bronchit)
  points <- cmdscale(1-bronchit.rf$proximity)
  xyplot(points[,2] ~ points[,1],
          groups=bronchit$rfac,
          xlab="Axis 1", ylab="Axis 2",
          par.settings=parset, aspect=1,
          auto.key=list(columns=2))
}

```

```

fig9.16 <- function(){
  opar <- par(mar=c(2,3,2,5))
  plot(aupts, axes=FALSE, ann=FALSE, fg="gray",
       frame.plot=TRUE)
  city <- rownames(aupts)
  pos <- rep(1,length(city))
  pos[city=="Melbourne"]<- 3
  pos[city=="Canberra"] <- 4
  pos[city=="Sydney"] <- 4
  par(xpd=TRUE)
  text(aupts, labels=city, pos=pos)
  par(opar)
}

```

Figure 1: Resistance in ohms is plotted against apparent juice content. A smooth curve (in gray) has been added, using the `lowess` smoother. The width of the smoothing window was the default fraction $f = \frac{2}{3}$ of the range of values of the x -variable.

Figure 2: Estimated contours of left eye responses to visual stimulæ, projected onto the plane.

```
figset9 <- function(){
  library(MASS)
  library(lattice)
  library(DAAGviz)
  if(!require('DAAG', quietly=TRUE))stop('DAAG must be installed')
  if(!requireNamespace('mgcv', quietly=TRUE))stop('mgcv must be installed')
  if(!requireNamespace('oz', quietly=TRUE))stop('oz must be installed')
  if(!requireNamespace('ggplot2', quietly=TRUE))
    stop('ggplot2 must be installed')
}
```

```
figset9()
```

```
Attaching package: 'DAAG'
The following object is masked from 'package:MASS':
  hills
```

```
## ---- bronchit-rfac ----
## Now make the outcome variable a factor
bronchit <-
  within(bronchit,
    rfac <- factor(r, labels=c("abs","pres")))
```

```
## ---- getErie ----
Erie <- greatLakes[, "Erie"]
```

```
library(rpart, quietly=TRUE)
```

Figure 3: Yields from 4 packages of land on each of eight sites on the Caribbean island of Antigua. Data are a summarized version of a subset of data given in Andrews and Herzberg 1985, pp.339-353.

Figure 4: Level of Lake Erie, in meters.

Figure 5: Panel A plots Lake Erie levels vs levels at lags 1, 2 and 3 respectively. Panel B shows a consistent pattern of decreasing autocorrelation at successive lags.

Figure 6: GAM smoothing term, fitted to the Lake Erie Data. Most of the autocorrelation structure has been removed, leaving residuals that are very nearly independent.

Figure 7: The plots are from repeated simulations of an AR1 process with a lag 1 correlation of 0.85. Smooth curves, assuming independent errors, have been fitted.

Figure 8: Predictions, 15 years into the future, of lake levels (m). The shaded areas give 80% and 95% confidence bounds.

Figure 9: Estimated contributions of the model terms to `mdbRain`, in a GAM model that is the sum of smooth terms in `Year` and `Rain`. The dashed curves show pointwise 2-SE limits, for the fitted curve. Note the downturn in the trend of `mdbRain` after about 1985,

Figure 10: The top left panel shows the autocorrelation plot of the residuals from the GAM model `mdbRain.gam`. The five remaining panels show autocorrelation plots for a series of independent random normal numbers.

Figure 11: Estimated number of events (aircraft crashes) per time interval versus time. In Panel A, the outcome variable was events per day, while in Panel B it was events per week.

Figure 12: Visual representation of a classification rule, derived using *linear discriminant analysis*, for the forensic glass data. A five-dimensional pattern of separation between the categories has been collapsed down to two dimensions. Some categories may therefore be better distinguished than is evident from this figure.

```
## ---- aupoints ----  
aupts <- cmdscale(audists)
```

Figure 13: Decision tree for predicting whether a miner has bronchitis. Where the condition at a node is satisfied, the left branch is taken. Thus, at the initial node, `cig<4.385` takes the branch to the left. In general, unless a random number seed is specified, the tree may be different for each different run of the calculations.

Figure 14: Each tree is for a different bootstrap sample of observations. The final classification is determined by a random vote over all trees. Where there are > 2 explanatory variables (but not here) a different random sample of variables is typically used for each different split. The final classification is determined by a random vote over all trees.

Figure 15: The plot is designed to represent, in two dimensions, the random forest result. It aims to reflect probabilities of group membership given by the analysis. It is not derived by a 'scaling' of the feature space.

Figure 16: Relative locations of Australian cities, derived from road map distances, using metric scaling.

Figure 17: In Panel A, Figure 16 has been rotated and scaled, to give a best fit to a map of Australia. Each city moves as shown by the line that radiates out from it. Panel B is the equivalent plot for the Sammon scaling ordination.